



E E S S I

EUROPEAN ENVIRONMENT FOR
SCIENTIFIC SOFTWARE INSTALLATIONS

Extending Arm's Reach by Going EESSI

ISC'24 - Hamburg - 15 May 2024

Kenneth Hoste (Ghent University, Belgium)

kenneth.hoste@ugent.be

Kenneth Hoste: Computer scientist from Belgium



- HPC system administrator + user support at Ghent University since Oct 2010
- BDFL of **EasyBuild** - tool to build & install **scientific software on HPC clusters**
- Active contributor to **EESSI**, partner in MultiXscale EuroHPC CoE
- Fan of open source software (FOSS), beers, and stickers
- [“How To Make Package Managers Cry” talk](#) at FOSDEM'18 (if you haven't seen it, you should!)

Scientific software is a different breed

- Scientists may not be trained software engineers (and that's OK)
- They often need help to get the software installed properly (especially on HPC systems)
- **Standard packaging tools do not suffice**, we ~~want~~ need to **build from source**
- Scientific software should be compiled for system on which it will be used (performance!)

I tried running
`sudo apt-get install pytorch`
on the CentOS cluster,
but it didn't work, even after I
entered my password correctly!



brilliant scientist




HPC user support

Optimized scientific software installations


- Software should be optimized for the system it will run on
- Impact on performance is often significant for scientific software

- Example: GROMACS 2020.1 (PRACE benchmark, Test Case B)
- Metric: (simulated) ns/day, higher is better
- Test system: dual-socket Intel Xeon Gold 6420 (Cascade Lake, 2x18 cores)
- **Performance of different GROMACS binaries, on exact same hardware/OS**



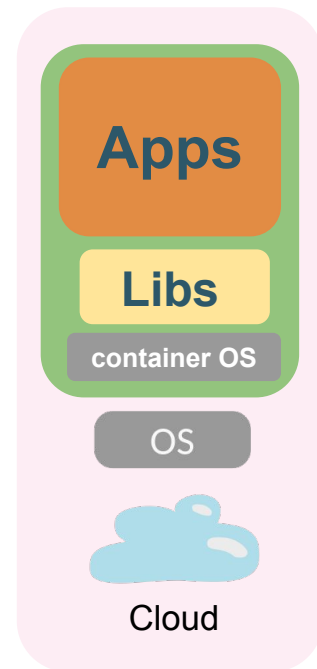
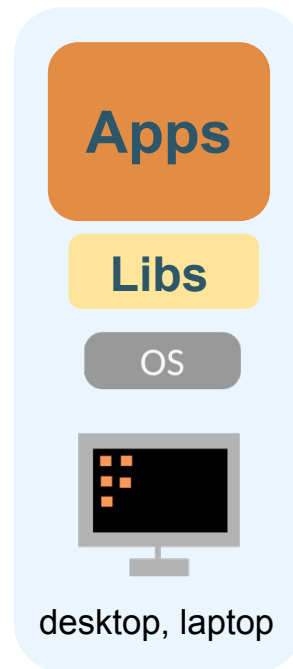
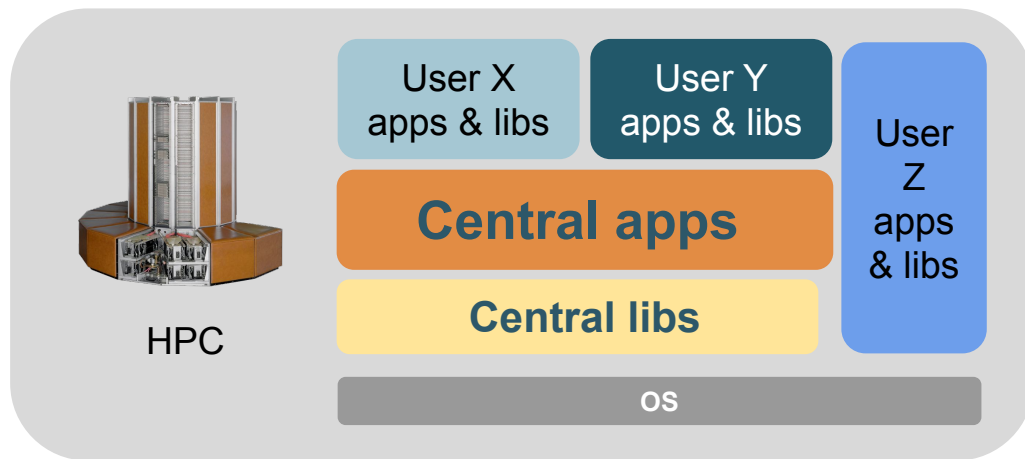
Landscape of scientific computing is changing

- **Explosion of available scientific software** applications (bioinformatics, AI, ...)
- Increasing interest in **cloud** for scientific computing (flexibility!)
- **Increasing variety in processor (micro)architectures** beyond Intel & AMD:
Arm is coming already here (see [Fugaku](#), [JUPITER](#), ...), RISC-V is coming (soon?)
- Broader adoption of **accelerated computing**, beyond NVIDIA GPUs
- In strong contrast: available (wo)manpower in **HPC support teams is (still) limited...**



We need to collaborate more...

- **Too much software** for a single support team to handle
- Different systems (CPU, GPU, OS, ...) => different problems
- Existing tools (EasyBuild, Spack) are **not sufficient anymore...**
- Lots of **duplicate work** across HPC sites and scientists
- **Diverse software stacks** across different platforms



European Environment for Scientific Software Installations

- **Public repository of (optimized!) scientific software installations**
- Avoid duplicate work by collaborating on a shared software stack
- Uniform way of providing software to users, regardless of the system they use!
- Should work on any Linux OS (incl. WSL) and system architecture
 - From laptops and personal workstations to HPC clusters and cloud
 - Support for different CPUs (AMD, Intel, Arm, RISC-V), interconnects, GPUs, etc.
- **Focus on performance, automation, testing, collaboration**






E E S S I


EUROPEAN ENVIRONMENT FOR
SCIENTIFIC SOFTWARE INSTALLATIONS

<https://eessi.io>

<https://eessi.io/docs>




CernVM-FS



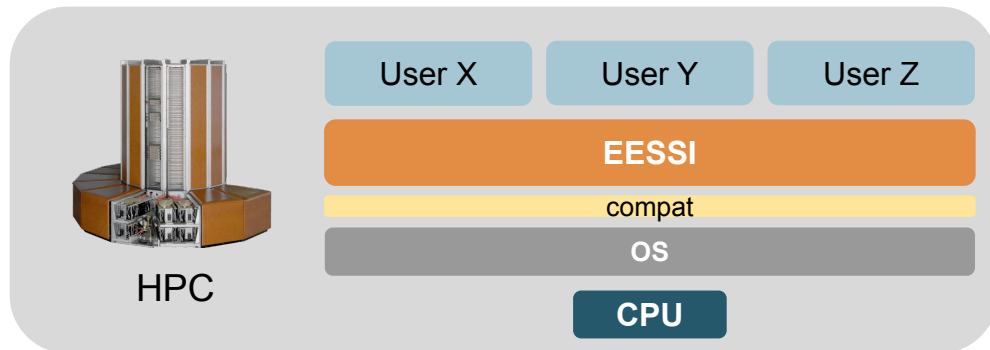
E E S S I
EUROPEAN ENVIRONMENT FOR
SCIENTIFIC SOFTWARE INSTALLATIONS

EESSI to the rescue!



E E S S I

EUROPEAN ENVIRONMENT FOR
SCIENTIFIC SOFTWARE INSTALLATIONS




**Shared repository of
(optimized) scientific
software installations**

**Same software stack
everywhere!**

eessi.io

eessi.io/docs

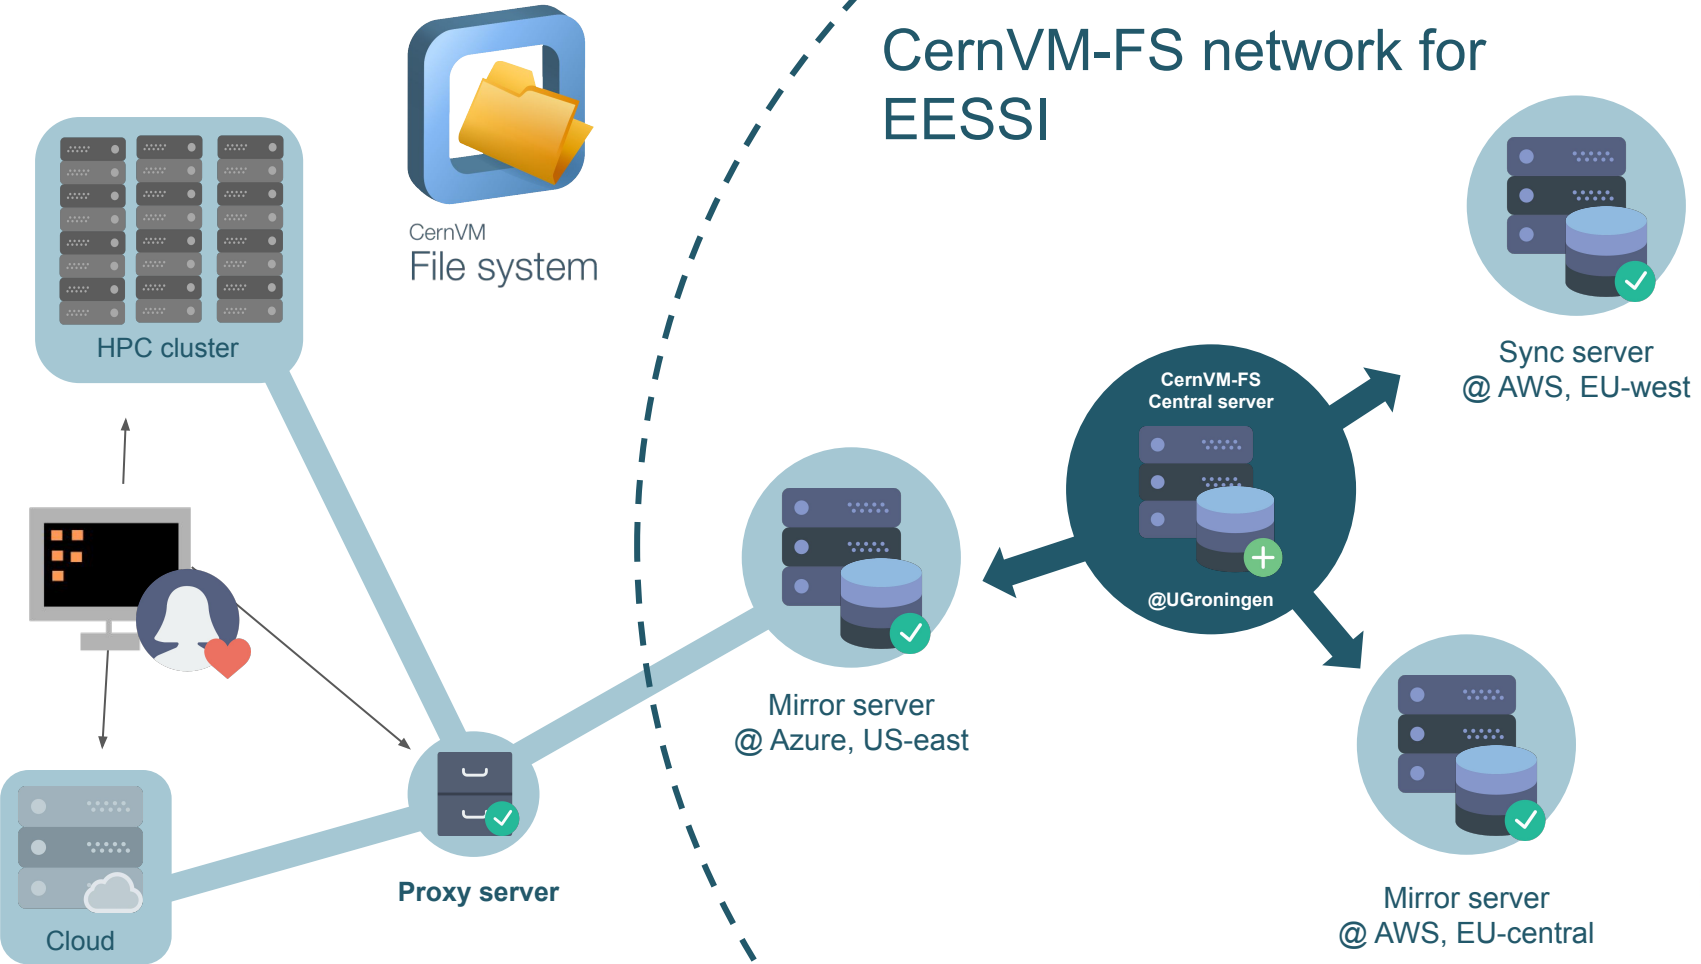
eessi.io/docs/support



EESSI ingredients: filesystem layer




CernVM-FS network for EESSI




Tutorial “Best Practices for CernVM-FS in HPC”



- <https://multixscale.github.io/cvmfs-tutorial-hpc-best-practices>
- Held online on 4 Dec 2023 (~3 hours), **recorded & available on YouTube**
- Over 200 registrations, ~125 attending the meeting
- Lecture + hands-on demos
- Topics:
 - Introduction to CernVM-FS + EESSI
 - Configuring CernVM-FS: client, Stratum 1 mirror server, proxy server
 - Troubleshooting problems
 - Benchmarking of start-up performance

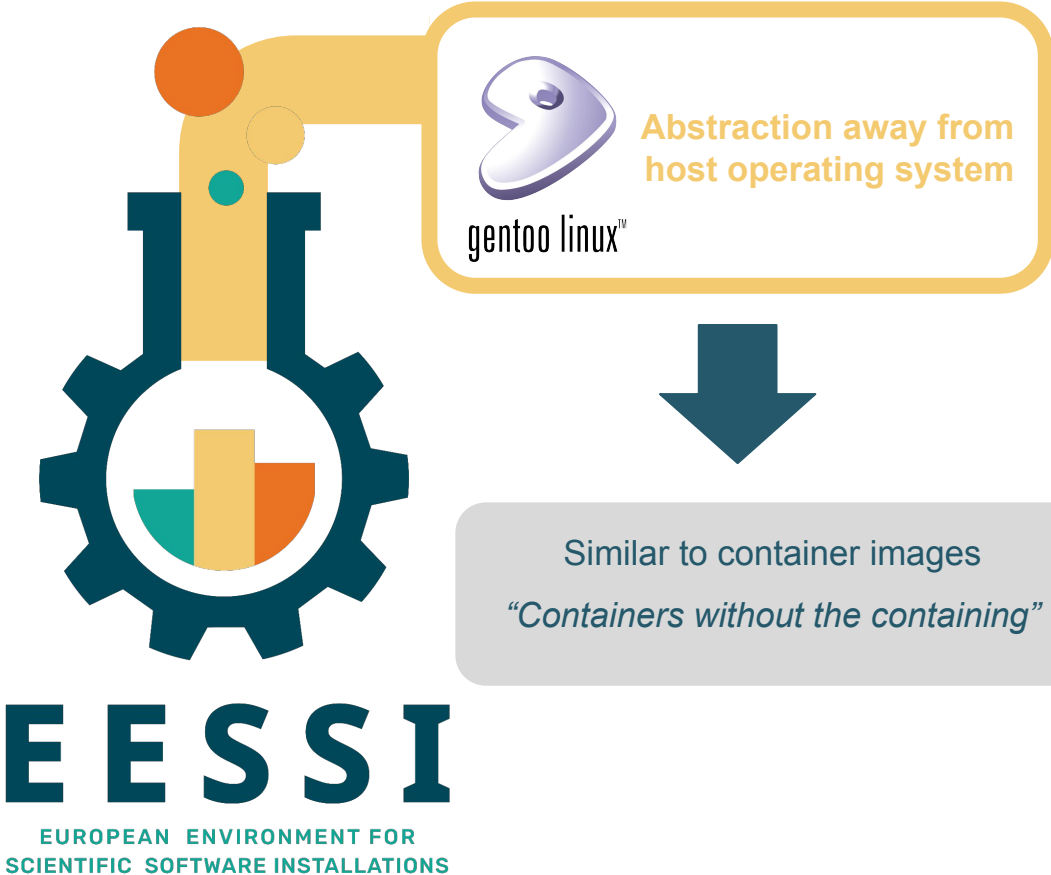


EESSI ingredients: compatibility layer



Global distribution of software installations

CernVM
File system




EESSI compatibility layer

github.com/EESSI/compatibility-layer



- “*Containers without the containing*”
- **Minimal collection of tools and libraries** (incl. glibc, bash, Python, Lmod, ...)
- **Built from source per CPU family** (x86_64, aarch64, ...) with [Gentoo Prefix](#)
- Installations included in software layer **only link to compat layer** (RPATH)
- Ensures **compatibility** with any client system running Linux

```
$ ls /cvmfs/software.eessi.io/versions/2023.06/compat/linux/aarch64/  
bin etc lib lib64 opt reprod run sbin stage1.log stage2.log  
stage3.log startprefix tmp usr var
```



EESSI ingredients: software layer



gentoo linux™

Abstraction away from
host operating system



CernVM
File system

Global distribution of
software installations



E E S S I


EUROPEAN ENVIRONMENT FOR
SCIENTIFIC SOFTWARE INSTALLATIONS



Optimized software
installations for specific
CPU microarchitectures

Intuitive user interface

module avail,
module load, ...




Automatic selection of
best suited part of
software stack for
host CPU microarchitecture

EasyBuild (in a nutshell)

<https://easybuild.io>



- EasyBuild is a tool to make **installing (scientific) software on HPC systems** easier
- **Building from source** + optimizing for specific CPU is strongly preferred (performance!)
- Created in 2009 by HPC-UGent team, open source (GPLv2) since 2012
- Now a **world-wide community** (> 400 unique contributors, close to 3,000 PRs per year)
- ~3,500 different software projects supported (excl. versions) + ~2,000 “extensions”



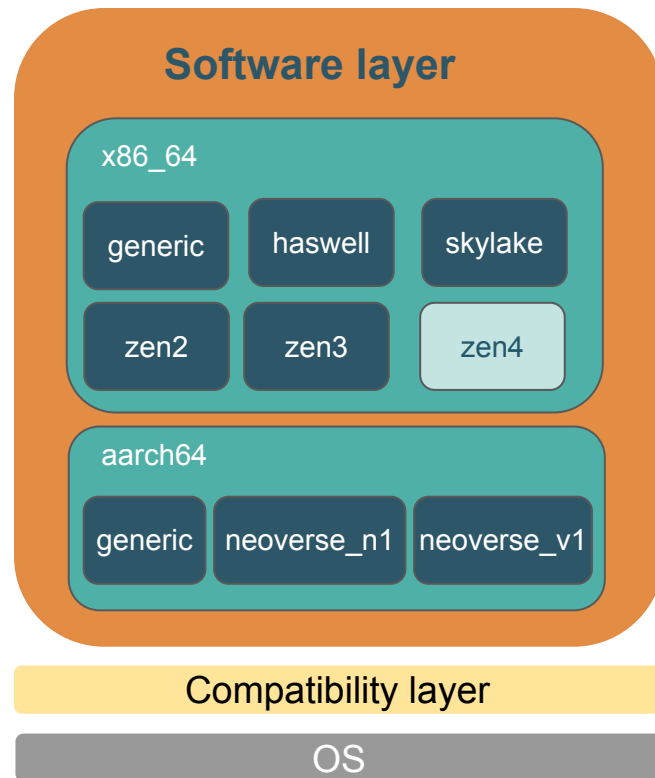

Software layer

- Installations of scientific software applications
- **Optimized for specific CPU targets**
- Works on any client system running Linux, since we only link to libraries in compat layer
- Built using [EasyBuild](#)
- Environment modules as user interface (via [Lmod](#))
- Detection of host CPU via [archspeg](#) (Python) or archdetect (bash)
- **Best subset of software installations for host CPU is automatically selected**

github.com/EESSI/software-layer




Lmod



Overview of installed software



- ~600 software software installations available per CPU target, and increasing every day
 - Including ESPResSo, GROMACS, LAMMPS, OpenFOAM, PyTorch, R, QuantumESPRESSO, TensorFlow, WRF
 - eessi.io/docs/available_software (coming soon!)
- Includes (**but is not limited to!**) applications specific to the MultiXscale CoE
- Focus on recent compiler toolchains: currently targeting foss/2023a and foss/2023b



Getting access to EESSI



- Native installation of CernVM-FS (*requires admin privileges*)
eessi.io/docs/getting_access/native_installation
- Using a container (via Apptainer)
eessi.io/docs/getting_access/eessi_container
- Via cvmfsexec
github.com/cvmfs/cvmfsexec

To check whether you have access to EESSI:

```
ls /cvmfs/software.eessi.io
```

Accessing EESSI via CernVM-FS



```
# Native installation
# Installation commands for RHEL-based distros
# like CentOS, Rocky Linux, AlmaLinux, Fedora, ...

# install CernVM-FS
sudo yum install -y
https://ecsft.cern.ch/dist/cvmfs/cvmfs-release/cvmfs-release-latest.noarch.rpm
sudo yum install -y cvmfs

# create client configuration file for CernVM-FS
# (no proxy, 10GB local CernVM-FS client cache)
sudo bash -c "echo 'CVMFS_CLIENT_PROFILE='single'' > /etc/cvmfs/default.local"
sudo bash -c "echo 'CVMFS_QUOTA_LIMIT=10000' >> /etc/cvmfs/default.local"

# Make sure that EESSI CernVM-FS repository is accessible
sudo cvmfs_config setup
```

Alternative ways of accessing EESSI are available, via a container image, via `cvmfsexec`, ...
eessi.io/docs/getting_access/native_installation - eessi.io/docs/getting_access/eessi_container

Using EESSI

eessi.io/docs/using_eessi/eessi_demos



```
/cvmfs/software.eessi.io/versions/2023.06/software
```

```
`-- linux
  |-- aarch64
  |   |-- generic
  |   |-- neoverse_n1
  |   |-- neoverse_v1
  |       |-- modules
  |       |-- software
  |-- x86_64
  |   |-- amd
  |   |   |-- zen2
  |   |   |-- zen3
  |   |   |-- zen4
  |   |-- generic
  |-- intel
  |   |-- haswell
  |   |-- skylake_avx512
```

```
$ source /cvmfs/software.eessi.io/versions/2023.06/init/bash
Found EESSI pilot repo @
/cvmfs/software.eessi.io/versions/2023.06!
```

```
archdetect says aarch64/neoverse_v1
Using aarch64/neoverse_v1 as software subdirectory
```

```
...
Environment set up to use EESSI pilot software stack, have fun!
```

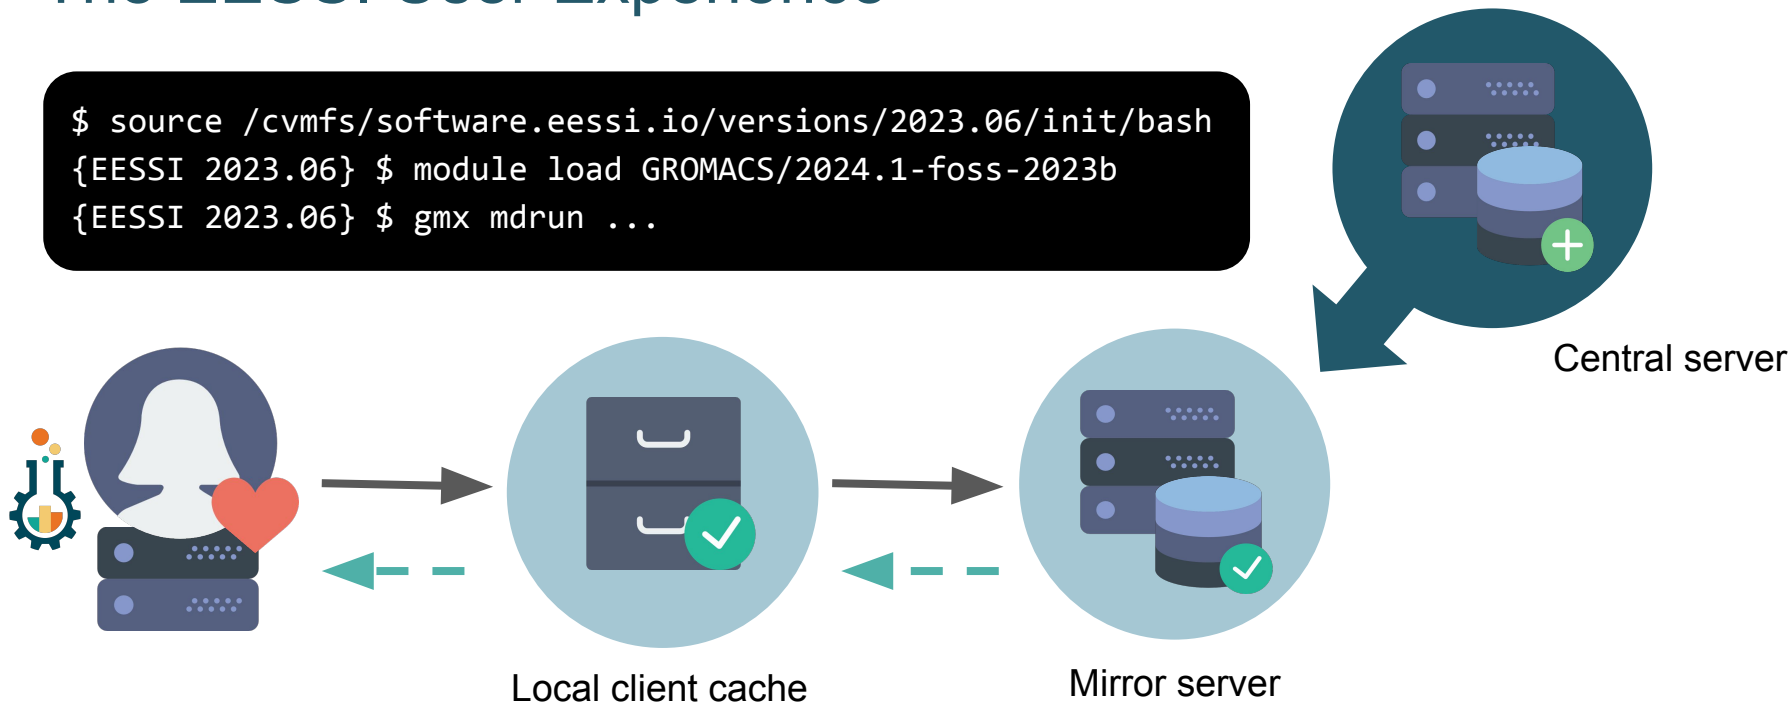
```
{EESSI 2023.06} $ module load R/4.3.2-gfbb-2023a
```

```
{EESSI 2023.06} $ which R
/cvmfs/software.eessi.io/versions/2023.06/software/linux/aarch64
/neoverse_v1/software/R/4.3.2-gfbb-2023a/bin/R
```

```
{EESSI 2023.06} $ R --version
R version 4.3.2
```

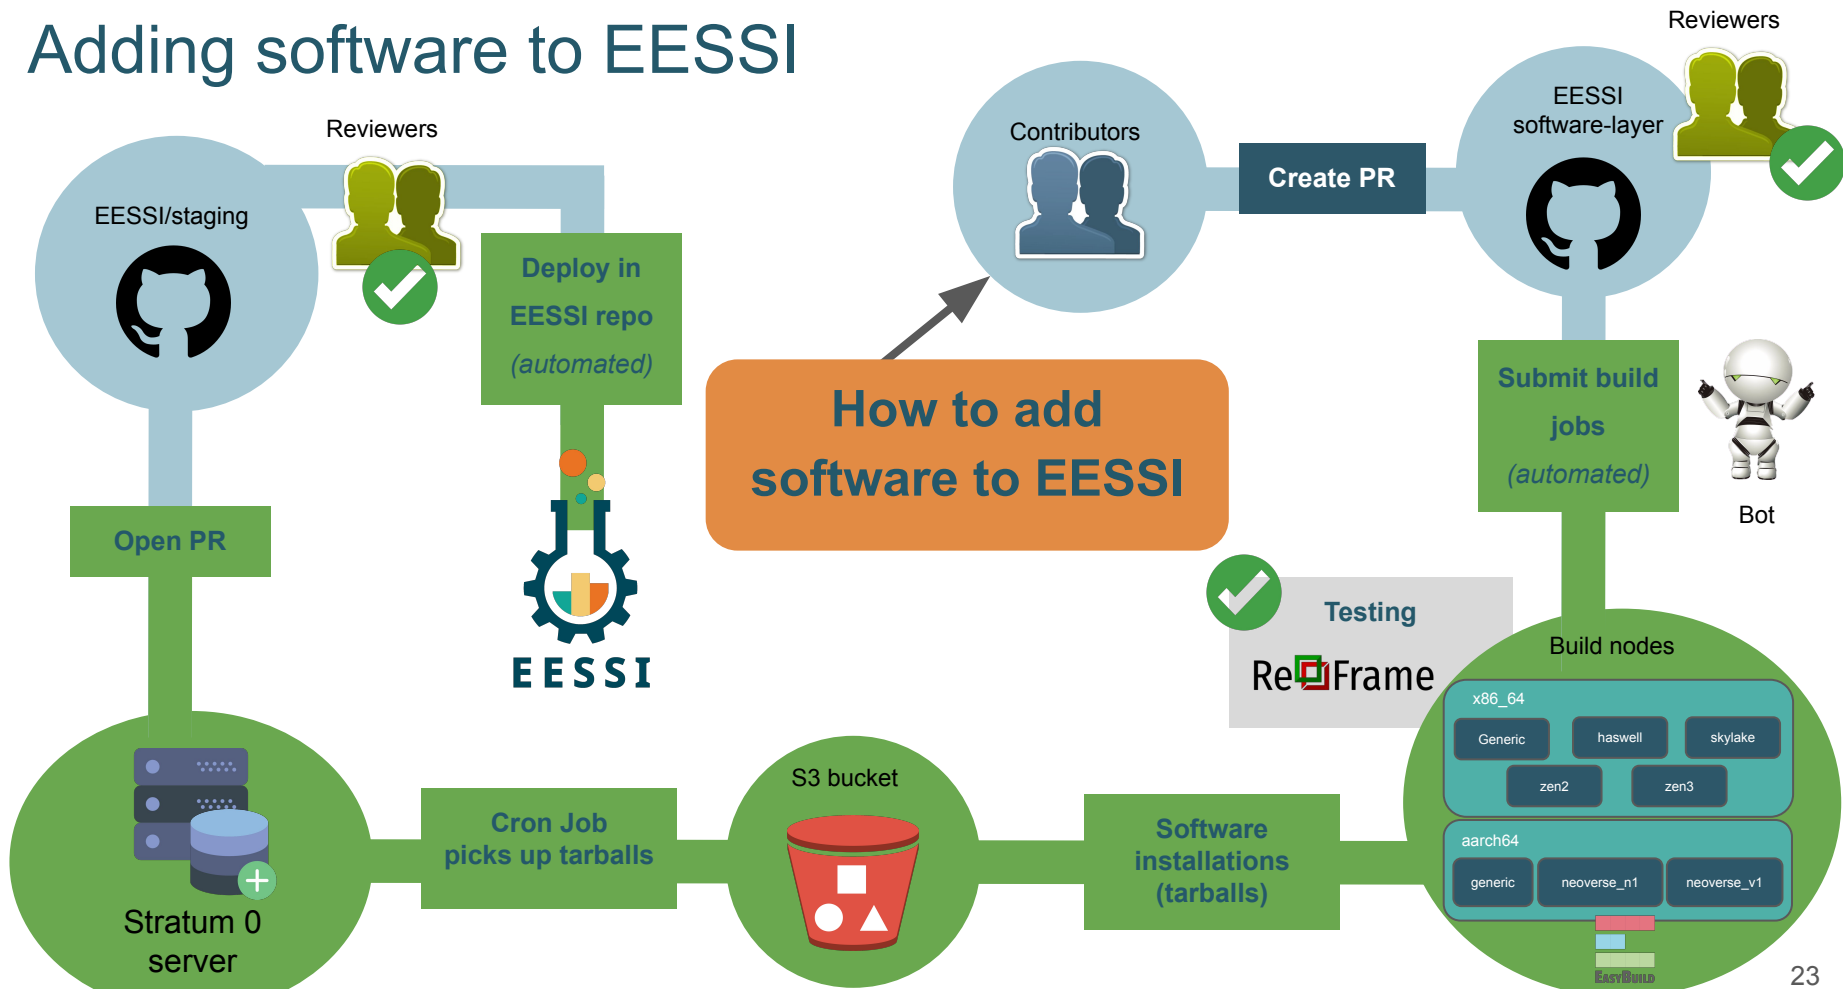
The EESSI User Experience

```
$ source /cvmfs/software.eessi.io/versions/2023.06/init/bash  
{EESSI 2023.06} $ module load GROMACS/2024.1-foss-2023b  
{EESSI 2023.06} $ gmx mdrun ...
```



EESSI provides **on-demand streaming**
of (scientific) software (like music, TV-series, ...)

Adding software to EESSI



EESSI test suite

ReFrame

eessi.io/docs/test-suite



- Ensure quality of the software installations provided by EESSI
- A suite of **portable tests** for scientific software, using ReFrame
- Version 0.2.0 includes tests for GROMACS, TensorFlow, OSU Micro Benchmarks
- WIP: tests for LAMMPS, QuantumESPRESSO, CP2K, OpenFOAM, ESPResSo, PyTorch
- See also MultiXscale deliverable D1.2 “**Plan for the design of a portable test suite**”

<https://zenodo.org/records/10451718>

Problems we have found through software testing (1/2)



Hang/crash in Open MPI's `smcuda` Byte Transport Layer (BTL) component

- See <https://gitlab.com/eessi/support/-/issues/41>
- Upstream issue: <https://github.com/open-mpi/ompi/issues/12270>
- Causes hanging or failing tests for FFTW, OpenFOAM, ESPResSo, ...
- Only happens on Arm Neoverse V1 (AWS Graviton 3)
- Fixed by Luke Robinson (AWS), see <https://github.com/open-mpi/ompi/pull/12338>
- Fix will be included in upcoming Open MPI release (v4.1.7+)
- Patch already applied in Open MPI installations included in EESSI 2023.06

Problems we have found through software testing (2/2)



Failing tests in GROMACS test suite

- See <https://gitlab.com/eessi/support/-/issues/47>
- Filesystem race when running tests concurrently ([GROMACS PR #4066](#))
- **Bug in SVE support, leading to (very) wrong results for several tests**
 - See <https://gitlab.com/gromacs/gromacs/-/issues/5057>
 - Works fine on A64FX (512-bit SVE), but **problem on Neoverse V1 (256-bit SVE)**
 - WIP fix in https://gitlab.com/gromacs/gromacs/-/merge_requests/4299
 - Will be fixed in upcoming GROMACS release (2024.2?)
 - Workaround for GROMACS 2024.1 in EESSI: use `-DGMX_SIMD=ARM_NEON_ASIMD`

Getting support for EESSI

- Via GitLab, or via email: support@eessi.io
- Report problems
- Ask questions
- Request software
- Get help with contributing
- Suggest features
- Confidential tickets possible (security issues, ...)



gitlab.com/eessi/support

Q Search or go to...

EESSI / EESSI support portal


Project

- EESSI support portal
- Manage >
- Plan >
- Code >
- Build >
- Deploy >
- Operate >
- Monitor >
- Analyze >

Help

README.md

EESSI support portal

MultiScale  **EESSI**
EUROPEAN ENVIRONMENT FOR SCIENTIFIC SOFTWARE INSTALLATIONS

Thanks to the [MultiScale EuroHPC project](#) we are able to provide support to the u

Contact

Create an issue with you GitLab account




If you have a GitLab account or create one you can create and manage your issue - also use one of our issue templates.

Contact us via E-mail

If you do not have a GitLab account you can also ask for support via E-mail.

Dedicated support team, thanks to EuroHPC Centre-of-Excellence




- EuroHPC Centre of Excellence: 4 year project (2023-2026), €6M budget (50% for EESSI)
- Collaboration between EESSI and CECAM: total of 16 partners (academic + industry)
- EESSI focuses on technical aspects: providing a shared stack of scientific software
- Scientific target: multiscale simulations with 3 key use cases
 - Helicopter design and certification for civil transport 
 - Battery applications to support the sustainable energy transition 
 - Ultrasound for non-invasive diagnostics and biomedical applications 

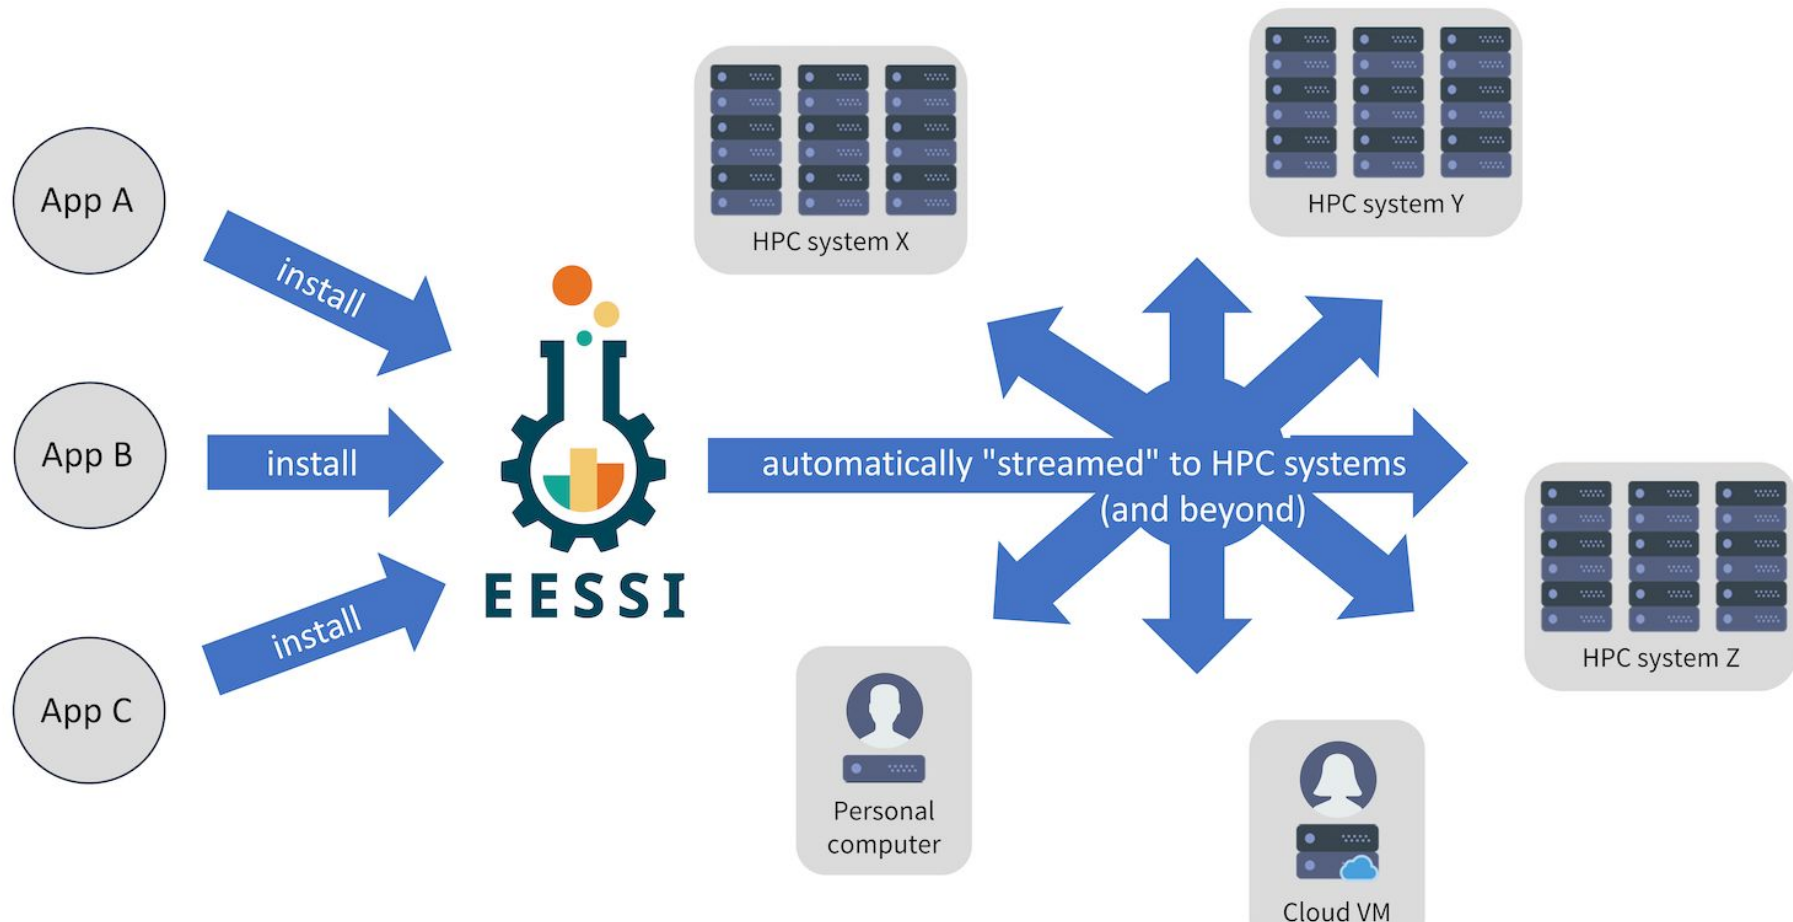
The role of EESSI in **Multi****scale**



automatically "streamed" to EuroHPC systems
(and beyond)



EESSI as a shared software stack



EESSI in a nutshell

- **On-demand streaming of optimized** scientific software installations
- **Works on any Linux distribution** thanks to EESSI compat layer
- **Uniform software stack** across various systems: laptop, HPC, cloud, ...
- Community-oriented: **let's tackle the challenges we see together!**



E E S S I

EUROPEAN ENVIRONMENT FOR
SCIENTIFIC SOFTWARE INSTALLATIONS



Acknowledgements



Co-funded by
the European Union



EuroHPC
Joint Undertaking

- Funded by the European Union. This work has received funding from the European High Performance Computing Joint Undertaking (JU) and countries participating in the project under grant agreement No 101093169.



- Thanks to Amazon Web Services (AWS) and Microsoft Azure for generously sponsoring the EESSI project with cloud credits, feedback, and guidance.





Website: eessi.io

GitHub: github.com/eessi

Documentation: eessi.io/docs

YouTube channel: youtube.com/@eessi_community

Paper (open access): doi.org/10.1002/spe.3075

EESSI support portal: gitlab.com/eessi/support

[Monthly online meetings](#) (first Thursday, 2pm CEST)

Join our mailing list & Slack channel